

Experimental Design and Data Analysis, Lecture 3

Eduard Belitser

VU Amsterdam

Lecture overview

- ① two paired samples (normal and not normal)
 - permutation test
 - dependence in two paired samples
 - Pearson's correlation test
 - Spearman's rank correlation test
- ② two independent samples (normal and not normal)
 - two samples t -test
 - Mann-Whitney test
 - Kolmogorov-Smirnov test

permutation tests for two paired samples

t-test is based on normality

sign test is for median

W - sign test is for median

-> permutation test is for general non-normality

Reminder: setting and design for two paired samples

Setting:

- An experiment with a **numerical outcome** measured according to **two conditions** per experimental unit;
- Interest is in a possible **difference** between the two outcomes per unit.

EXAMPLE Difference in **average course grade** for **mathematical courses** and **informatics courses** for BA-students at the VU.

EXAMPLE Difference in **pain relief** by an **active drug** and a **placebo** for patients.

Design (the standard paired samples design):

- Take a random sample of experimental units from the relevant population.
- Measure the two outcomes on each unit.



Idea of permutation technique

- Data $(X_1, Y_1), (X_2, Y_2), \dots, (X_N, Y_N)$.
- In a **permutation test** **we do not assume normality**.
- We can use **any test statistic** $T = T(X_1, Y_1, \dots, X_N, Y_N)$ to test the null hypothesis of no difference between the distribution of X_i and that of Y_i within samples. The choice depends on the difference conjectured.
- Like in a bootstrap test, we simulate the distribution of T under H_0 , using B surrogate T^* -values. Repeat B times (for $i = 1, \dots, B$):
 - generate (X_j^*, Y_j^*) by generating a **permutation** of the original (X_j, Y_j) (relabeling) for $j = 1, \dots, N$, i.e., choose between (X_j, Y_j) and (Y_j, X_j) with equal probability.
 - compute $T_i^* = T(X_1^*, Y_1^*, \dots, X_N^*, Y_N^*)$
- Under H_0 of no difference between the distributions of X and Y within pairs permuting the labels does not change the distribution of T .

Permutation test in R: data input and graphics

Recall dataset `ashina.txt` (drug or placebo against headache for 16 subjects).

```
> ashina=read.table("ashina.txt",header=TRUE)
> boxplot(ashina[,1],ashina[,2],names=c("active","placebo"))
> plot(ashina[,1],ashina[,2]); abline(0,1)
```



(Based on this picture we expect the active medicine to yield better pain relief.)

Permutation test in R: testing (1)


```
> mystat=function(x,y) {mean(x-y)}
> B=1000
> tstar=numeric(B)
> for (i in 1:B)
+ {
+   ashinastar=t(apply(cbind(ashina[,1],ashina[,2]),1,sample))
+   tstar[i]=mystat(ashinastar[,1],ashinastar[,2])
+ }
> myt=mystat(ashina[,1],ashina[,2])
```

Instead of computing all $2^{16} = 65536$ possible permutations, we generate 1000 randomly chosen permutations to estimate the distribution of our test statistic under H_0 . The function `apply` applies a function to either all rows or all columns in a matrix.

Permutation test in R: testing (2)

```
> myt
[1] -42.875
> hist(tstar)
> pl=sum(tstar<myt)/B
> pr=sum(tstar>myt)/B
> p=2*min(pl,pr)
> p
[1] 0.008
```

Conclusion: there is indeed a significant difference between the active drug and the placebo.



Permutation test: discussion

- A permutation test for two paired samples can be performed with **any test statistic** that **expresses difference between the X and Y within pairs.** (The mean of differences $Z_i = X_i - Y_i$ is most common to consider, but one may as well consider the median of the Z_i 's. Then the test is a bootstrap version of the sign test on the median of Z_i equal to 0.)
- **Nonparametric alternatives** to the permutation test for two paired samples are the sign test and the Wilcoxon signed rank test applied to the differences (cf. the previous lecture).

Dependence in two paired samples

Dependence between two paired samples

Setting:

An experiment with two **numerical outcomes** (say X and Y) per experimental unit. Interest is in a possible **dependence** between the two outcomes per unit.

EXAMPLE Relation between **shoe size** and **body mass index** of a person.

EXAMPLE Relation between **average course grade** and **number of students taking the course** for courses at the VU.

EXAMPLE Relation between amount of **precipitation** and **sun hours** for different cities in Europe.

Design:

- Take a random sample of experimental units from the relevant population.
- Measure the two quantities on each unit. (The two outcomes are in principal related, because measured on the same experimental unit.)
- However, we possibly have measured unrelated quantities of the units and we want to test whether these quantities are **correlated**.

Pearson's correlation test

- Data $(X_1, Y_1), (X_2, Y_2), \dots, (X_N, Y_N)$.
- The **Pearson correlation test** assumes **normality** of the both X_i 's and Y_i 's.
(Rather, the asympt. normality of the sample correlation $\hat{\rho}$.)
- The test is based on the sample correlation coefficient (which **estimates** of the “true” correlation $\rho = \text{cor}(X, Y)$):

$$\hat{\rho} = \hat{\rho}_{X,Y} = \frac{\sum_{i=1}^N (X_i - \bar{X}_N)(Y_i - \bar{Y}_N)}{\sqrt{\sum_{i=1}^N (X_i - \bar{X}_N)^2 \sum_{i=1}^N (Y_i - \bar{Y}_N)^2}}.$$

- We **test** the null hypothesis $H_0 : \rho = \rho_0 = 0$ that the correlation between the two populations is $\rho_0 = 0$. The **test statistic** is given by

$$T_{\rho} = \frac{\hat{\rho} - \rho_0}{\left(\frac{1-\hat{\rho}^2}{n-2}\right)^{1/2}} = \frac{\hat{\rho}}{\left(\frac{1-\hat{\rho}^2}{n-2}\right)^{1/2}},$$

which has under $H_0 : \rho = 0$ a **t-distribution** with $n - 2$ degrees of freedom.

Spearman's rank correlation test

- Data $(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)$.
- **Spearman's rank correlation test** does **not assume normality**. The test considers the ranks $R(X_i)$ and $R(Y_i)$ in the two samples, and compares the ordering of the ranks in the X_i and the Y_i .
- If the data are rank correlated, these sequences of ranks will run (approximately) in parallel or in opposite order.
- The test statistic is the sample correlation $\tilde{\rho}$ between the **the rank vectors**.
- We **test** the null hypothesis $H_0 : \rho_s = 0$. (Correlation of the rank variables)

If all n ranks are distinct integers, the test statistic can be computed as

$\tilde{\rho} = 1 - \frac{6 \sum_{i=1}^n d_i^2}{n(n^2-1)}$ where $d_i = R(X_i) - R(Y_i)$ is the difference between the two ranks of observations X_i and Y_i .


Correlation tests in R: example

Consider the data frame `peruvians.txt`, where the rows correspond to 39 men that moved from a native culture to a modern society. Amongst others, weight and systolic blood pressure were measured.

```
> peruvians=read.table("peruvians.txt",header=TRUE); peruvians
  age migration weight length chin  arm calf wrist systolic diastolic
1   21          1  71.0  1629  8.0  7.0 12.7   88      170        76
2   22          6  56.5  1569  3.3  5.0  8.0   64      120        60
  [ some output deleted ]
39  54          40  87.0  1542 11.3 11.7 11.3   92      152        88
```

```
> attach(peruvians)
> plot(systolic~weight)
```

Based on this picture, we expect dependence between systolic and weight.



Pearson's test in R: example

```
> cor.test(systolic,weight)
```

Pearson's product-moment correlation

```
data:  systolic and weight
t = 3.7164, df = 37, p-value = 0.0006654
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.2463759 0.7186619
sample estimates:
      cor
0.5213643
```



Conclusion: there is significant correlation, if normality is assumed.

The default for `cor.test` is Pearson's correlation test, based on normality.

Correlation tests: diagnostics

Check the normality assumption on the two samples:

```
> qqnorm(weight,main="Q-Q Plot weight")  
> qqnorm(systolic,main="Q-Q Plot systolic")
```



QQ-plots show that normality is not plausible for the weight sample. Hence, use the rank correlation test of Spearman (and **not** Pearson's correlation test).

Spearman's test in R: example

```
> cor.test(systolic,weight,method="spearman")
```

Spearman's rank correlation rho

data: systolic and weight

S = 5322.352, p-value = 0.003119

alternative hypothesis: true rho is not equal to 0

sample estimates:

rho

0.4613004

Warning message:

In cor.test.default(systolic, weight, method = "spearman") :

Cannot compute exact p-values with ties

Conclusion: there is indeed significant rank correlation.

There is a warning about ties, which means that some values occur multiple times in weight and/or systolic. Therefore *R* uses an approximation for the *p*-value.

two independent samples

Two independent samples: setting and design

Setting: an experiment with

- one **numerical outcome** per experimental unit,
- two **groups** of experimental units.

Interest is in a possible **difference** between the two populations. medskip

EXAMPLE Comparing the **weight** of newborn children in **two countries**, The Netherlands and Chile.

EXAMPLE Measurement of **total yield** from an agricultural plot for **two different fertilizers**.

Design:

- Take a random sample of experimental units of size M from the first population and a random sample of size N from the second population;
- Measure the outcome on each unit.

The numbers M and N need not be the same. (Taking the number M and N equal is preferable since it maximizes the power of two sample tests.)

Analysis A: two samples t -test

- Data (X_1, \dots, X_M) and (Y_1, \dots, Y_N) .
- The **two samples t -test** assumes that both samples X_1, \dots, X_M and Y_1, \dots, Y_N come from a **normal** population. Denote the mean of the first population by μ and the mean of the second by ν .
- We **test** the null hypothesis $H_0 : \mu = \nu$ that the means of the populations are the same.
- The **test statistic** is


$$T = \frac{\bar{X}_M - \bar{Y}_N}{S_{N,M}}, \quad \text{which has the } t_{N+M-2}\text{-distribution under } H_0.$$

Analysis A in R — data input and graphics



Consider two data sets of measurements of the speed of light (minus 299000) by Michelson in 1879 and in 1882.

```
> light1=scan("light1.txt"); light2=scan("light2.txt")
> hist(light1); hist(light2); boxplot(light1,light2)
```

Histogram of light1



Histogram of light2



Analysis A in R — estimation and testing

The two samples t -test:

```
> t.test(light1,light2)
```

Welch Two Sample t-test



```
data: light1 and light2
t = 4.0598, df = 27.754, p-value = 0.0003625
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 47.63387 144.73135
sample estimates:
mean of x mean of y
 852.4000 756.2174
```

Conclusion: H_0 of equal means is rejected.

By default `t.test` with two arguments performs the two samples t -test for independent samples.

Analysis A in R — diagnostics

```
> qqnorm(light1)
> qqnorm(light2)
```



Normality of the second sample is actually doubtful.

Analysis B: the Mann-Whitney test

- Data (X_1, \dots, X_M) and (Y_1, \dots, Y_N) .
- The **Mann-Whitney test** assumes that the sample X_1, \dots, X_M stems from population F and sample Y_1, \dots, Y_N stems from population G .
- We **test** the null hypothesis $H_0 : F = G$ that the populations are the same.
- The Mann-Whitney test is again based on ranks. It considers the M ranks R_1, \dots, R_M of X_1, \dots, X_M in the combined sample $(X_1, \dots, X_M, Y_1, \dots, Y_N)$ of length $M + N$. If $F = G$ these M rank numbers should lie randomly between 1 and $M + N$. The test statistic is

$$T = \sum_{i=1}^M R_i, \quad \text{the distribution of } T \text{ under } H_0 \text{ is (approximately) known.}$$

- Large values of T indicate that F is shifted towards the right from G , i.e. that X -values are bigger than Y -values.

If responses are continuous, a significant result of Mann-Whitney test shows a difference in medians, actually this test is only consistent against the alternative $H_1 : P(X > Y) \neq P(Y > X)$ (or $P(X > Y) + 0.5P(X = Y) \neq 0.5$).

Analysis B in R — testing

```
> wilcox.test(light1,light2)
```

Wilcoxon rank sum test with continuity correction

data: light1 and light2

W = 1829, p-value = 1.056e-05

alternative hypothesis: true location shift is not equal to 0

Conclusion: H_0 of equal medians is rejected. The underlying distribution of light1 is shifted to the right from that of light2.




When given two arguments `wilcox.test` will perform the Mann-Whitney test for two samples. The Mann-Whitney test is especially suited for detecting shift differences — differences in location — between two populations.

Analysis C: Kolmogorov-Smirnov test

- Data (X_1, \dots, X_M) and (Y_1, \dots, Y_N) .
- The **Kolmogorov-Smirnov test** assumes that the sample X_1, \dots, X_M stems from population F and sample Y_1, \dots, Y_N stems from population G .
- We **test** the null hypothesis $H_0 : F = G$ that the populations are the same.
- The Kolmogorov-Smirnov test is based on the differences in the histograms of the two samples.
- The **test statistic** computes the maximal vertical difference in **summed histograms** (empirical distribution functions). Its distribution under H_0 is known (e.g., in R).

Analysis C in R — graphics

```
> hist(light1)
> hist(light2)
```



Analysis C in R — testing

```
> ks.test(light1,light2)
```

Two-sample Kolmogorov-Smirnov test

```
data: light1 and light2  
D = 0.5391, p-value = 3.803e-05  
alternative hypothesis: two-sided
```

Warning message:

```
In ks.test(light1, light2) : cannot compute exact p-values with ties
```

Conclusion: H_0 of equal means is rejected. The mean of light1 is larger.

There is a warning about ties again. *R* uses an approximation for computing the p -value.

To finish

Today we discussed: two samples tests; for paired and independent samples, and for normal and not normal cases.

Next time: k samples, one way ANOVA.